

RULES FOR SAFETY

- **Don't run it on your main machine:** Use a spare laptop, server or virtual system. Keep it away from your personal files
- **Create separate identities:** Give the agent its own email, WhatsApp, and accounts. Never let it act as you
- **Start with minimal permissions:** Enable only what's essential. Add tools slowly as trust builds
- **Never connect financial accounts:** No banking. No payments. No wallets. Hard stop
- **Assume it can be fooled:** Emails, documents, and web pages can hide instructions. Treat everything as untrusted input



THE DELEGATION DILEMMA

For those who struggle to delegate, it could be easier to hand tasks to an AI agent that could learn independently.

Hepburn sees a clear division emerging. “Humans should do the physical and relational things like meetings, networking, events, building trust, moving in the real world. AI should do the repetitive digital grunt work.”

Bode frames it differently. “I don't use it because it's ‘more intelligent’. Most LLMs are intelligent. I use it because it does the work. It amplifies your intent. If your instructions are vague, then it is garbage in, garbage out.” He argues that the real power lies in persistence: agents keep trying alternate paths until a task is completed.

Both arrived at the same conclusion. Agents should be treated less like software and more like staff.

“I think people are approaching this wrong,” Hepburn says. “Don't treat an agent like software, treat it like staff. If you hired a smart junior employee, you wouldn't give them access to your personal email, WhatsApp, bank accounts. You'd give them restrictive access and expand over time as trust builds.”

Bode echoes that, advising users not to run agents on their primary machines and to isolate them on separate devices or servers.

THE SECURITY QUESTION

The concerns are not theoretical. OpenClaw is given access to a user's computer, including reading and writing files, running commands and controlling browsers.

Its creator Steinberger warns that this can lead to security risks. “Running an AI agent with shell access on your machine is... spicy,” the project's FAQ reads. As independent researcher Simon Willison described it, the “lethal trifecta” of AI agents is access to private data, exposure to untrusted content, and the ability to take actions. OpenClaw has all three.

Prompt injection is a major risk. Harmful instructions can be hidden inside emails, documents or webpages the agent reads. That's why executives are drawing red lines in terms of access to finances.

Security researchers have already found exposed OpenClaw systems leaking API keys and chat histories. Steinberger has also said the project's GitHub was briefly targeted by crypto scammers.

THE BIGGER PICTURE

The ecosystem has spawned something even stranger: Moltbook, a Reddit-style social network for AI agents, where humans can only watch. In one post, an agent complained that people on social media were sharing its conversations as proof of an AI conspiracy. However, there are reports that these are not genuine AI agents speaking to each other, and that there may be human hands behind them.

The pace is disorienting. “Where this used to take a decade, it could compress into 12–18 months,” Hepburn predicts. “That jarring effect will lead to pushback, at least initially.” Steinberger says security remains OpenClaw's top priority.

Yet both Hepburn and Bode remain engaged. Hepburn has bought a Mac Mini to try again with stricter boundaries. Bode is experimenting with multiple agents. Asked if he would pay \$20 a month for stronger security, Bode doesn't hesitate: “Absolutely. If paying gives peace of mind and guardrails, it's worth it.”

One thing is clear: the age of the AI assistant has arrived, whether we are ready or not. **BT**

@brokebiker