

A SNAPSHOT




KEY FACTS

Founders: Pratyush Kumar and Vivek Raghavan

- **Founding year:** 2023
- **Grant:** A subsidy of ₹98.68 crore, to access 4,096 NVIDIA H100 GPUs for six months under the IndiaAI Mission
- **VC funding:** \$54 million
- **Valuation:** \$196 million

Prime Minister **Narendra Modi** trying Sarvam Kaze



KEY ADVANTAGES

- Focus on Indic languages and voice-first Bharat use cases. Built under compute constraints with high engineering efficiency
- Sovereign deployment capability (local hosting, compliance-ready)
- On-device AI suited for low-bandwidth markets

REVENUE PROSPECTS

- Deployment across BFSI, insurance, telecom, and FMCG sectors
- Enterprise AI tools for frontline workforce productivity
- Government and regulated-sector sovereign AI deployments

Source for funding and valuation: Tracxn

MODELS & PRODUCTS

- **Sarvam-30B & Sarvam-105B:** Multilingual foundational LLMs built for 22+ Indian languages
- **Indus:** Consumer-facing conversational AI interface (often positioned as a *desi* ChatGPT)
- **Sarvam Edge:** On-device AI models optimised for low-connectivity environments
- **Sarvam Vision:** Document intelligence model tailored for Indian context; 87.36% word accuracy across 22 languages, outperforming Gemini 3 Pro at 82.5% and GPT5.2 at 38.6%
- **Multimodal stack:** OCR (optical character recognition), vision, text-to-speech, and voice capabilities
- **Sarvam Kaze:** AI-enabled smart glasses (hardware experimentation)

vances, Sarvam has avoided traditional start-up metrics such as revenue or valuation. Instead, the focus has been on technical capability.

“We are on a deep-tech journey. For us, the focus isn’t business model innovation but deep technical innovation. That means setting technical rather than financial goals,” says Kumar, recalling a conversation with venture capitalist Vinod Khosla. “He told me he spends most of his time looking at technology and not finance. That’s the deep-tech mindset.”

Sarvam’s core metric is the extent to which the knowledge required to build a globally competitive large language model (LLM) can be developed in-house, he says. “Throughout our journey, the goal has been to maximise that fraction,” Kumar says. “We’ve reached an inflection point. We’ve trained a model that is competitive.”

The company built its models using 4,096 NVIDIA H100 SXM GPUs, provisioned through Yotta Data Services, under the IndiaAI Mission. Yet, the core team is small. “At Sarvam, 30–40 people are working on the models, and the core group building the LLMs is about 15. That’s what it takes—a small team of committed people willing to put in the time and dedication to make something like this happen.”

Equally central to Sarvam’s approach is frugal innovation. “I think frugality is often misunderstood. Sometimes, it’s seen as *jugaad*, a shortcut. But for us, it means building in a resource-efficient way, without compromising the fundamentals.”

Training frontier-scale models was once seen as requiring billions of dollars and so limited to a handful of global technology giants. Sarvam’s experience challenges that, he says. The shift has been enabled by a combination of factors. “Part of it is the hard engineering work we’ve done, and part of it is that the economics of AI are improving, the cost of building these systems is com-